

Time Optimal Reachability Analysis using Swarm Verification

Zhengkui Zhang, Brian Nielsen, Kim G. Larsen

Department of Computer Science, Aalborg University

{zhzhang,bnielsen,kgl}@cs.aau.dk

Contents


- Motivation
 - Time optimal reachability
 - Swarm verification
- Sequential Time Optimal Reachability
- Swarm Time Optimal Reachability
- Experiments
- Conclusion
- Future Work

Motivation – Time optimal reachability

- Methods for scheduling and planning problems:
 - Numerical: linear / dynamic programming, operation research, etc
 - Model based: time optimal reachability, timed game, etc
- Tools for model based method: Uppaal, Kronos.
- Advantages of model based method:
 - Model real-time behaviors, constraints in natural way;
 - Flexible choice of underlying algorithms efficiently implemented in model checker.

Motivation – Time optimal reachability

- Real time model checking: given a timed model \mathcal{M} and a specification ϕ , verify $\mathcal{M} \models \phi$, generate a diagnostic trace ρ if $\mathcal{M} \not\models \phi$.
- Time optimal reachability: given a timed model \mathcal{M} and a reachability specification ϕ , when $\mathcal{M} \models \phi$ generate a diagnostic trace ρ such that its **span** (accumulated delay) is **minimum**.
- Time optimal reachability is real time reachability problem with cost, optimality and pruning into account.



Contents

- Motivation
 - Time optimal reachability
 - Swarm verification
- Sequential Time Optimal Reachability
- Swarm Time Optimal Reachability
- Experiments
- Conclusion
- Future Work

Motivation – Swarm verification

- As the number of components in the model increases, **state explosion** problem appears.
- Two main paradigms to speedup exploration and alleviate state explosion problem:
 - Parallel- and distributed computing: DiVinE, LTSmin, etc
 - Swarm verification: Swarm-Spin
- Swarm verification: a large number of model checker instances run in parallel on a computer cluster, using different, typically randomized search strategies.


Motivation – Swarm verification

- Swarm verification weak points:
 - No data parallelism, barely cut down termination time;
 - Duplicate work, duplicate copy of (partial) state space.
- Swarm verification strong points:
 - Get high quality results fast before exploring full state space, especially under time- and memory constraints
 - Easy to implement.
- “Even though each instance is tiny, there is strength in numbers.”



Motivation – Swarm verification

- We propose swarm verification using Uppaal to accelerate time optimal reachability.
- Swarm instances can exchange better costs to make more efficient pruning on all instances.
- Four swarm algorithms:
 - P-RDFS,
 - S-RDFS,
 - S-Mix,
 - S-Agent.



Contents

- Motivation
- Sequential Time Optimal Reachability
- Swarm Time Optimal Reachability
- Experiments
- Conclusion
- Future Work

Sequential Time Optimal Reachability

Algorithm 1: Sequential Time Optimal Reachability

PASSED $\leftarrow \emptyset$, WAITING $\leftarrow \{(\ell_0, Z_0)\}$, COST $\leftarrow \infty$

ORDER $\in \{\text{DFS}, \text{BFS}, \text{RDFS}\}$

Procedure Main()

```
1  while WAITING  $\neq \emptyset$  do
2    Search()
3  return COST
```

Procedure Search()

```
4  select  $(\ell, Z)$  from WAITING on ORDER
```

```
5  if  $(\ell, Z) \models \text{Goal}$  then
```

```
6    if  $\text{MinCost}(\ell, Z) < \text{COST}$  then
```

```
7      COST  $\leftarrow \text{MinCost}(\ell, Z)$ 
```

```
8  else if  $(\ell, Z) \notin \text{PASSED}$  and  $\text{MinCost}(\ell, Z) < \text{COST}$  then
```

```
9    add  $(\ell, Z)$  to PASSED
```

```
10   forall the  $(\ell', Z')$  such that  $(\ell, Z) \rightsquigarrow (\ell', Z')$  do
```

```
11     add  $(\ell', Z')$  to WAITING
```

Maintain current best span




Update cost

Inclusion checking

Generate successors


Branch and bound elimination

Sequential Time Optimal Reachability



Sequential Time Optimal Reachability

- Branch and Bound (B&B)
 - Pruning technique, avoid enumerate full state space;
 - Skip nodes whose cost $>$ current best cost to the goal;
 - Heuristic function.



Contents

- Motivation
- Sequential Time Optimal Reachability
- **Swarm Time Optimal Reachability**
- Experiments
- Conclusion
- Future Work

Swarm Time Optimal Reachability

- Benefits of RDFS: random search, give results in inexhaustible state space.
 - Four swarm algorithms based on RDFS:
 - Primitive swarm: do **not exchange** cost
 - P-RDFS: all do RDFS, first-complete-terminate
 - Cooperative swarm: **exchange** cost
 - S-RDFS: as P-RDFS
 - S-Mix: **one (BFS)**, others (RDFS)
 - S-Agent: **root (BFS)**, **agents (RDFS)**, agents periodically ask states from Root to search from, root terminates agents.
- Is sharing cost effective?


Why BFS?

Swarm Time Optimal Reachability

- BFS is important in S-Mix and S-Agent:
 - In Uppaal, BFS runs much faster than DFS and RDFS.
 - BFS builds large zones while DFS/RDFS causes high level of segmentations.
- In a word, a BFS instance enables swarm verification terminate fast.

Swarm Time Optimal Reachability


- P-RDFS: primitive version of swarm



Algorithm in Appendix A

Swarm Time Optimal Reachability


- S-RDFS: Cooperative Swarm RDFS



Algorithm in Appendix B

Swarm Time Optimal Reachability

- S-Mix: Cooperative Swarm Mix

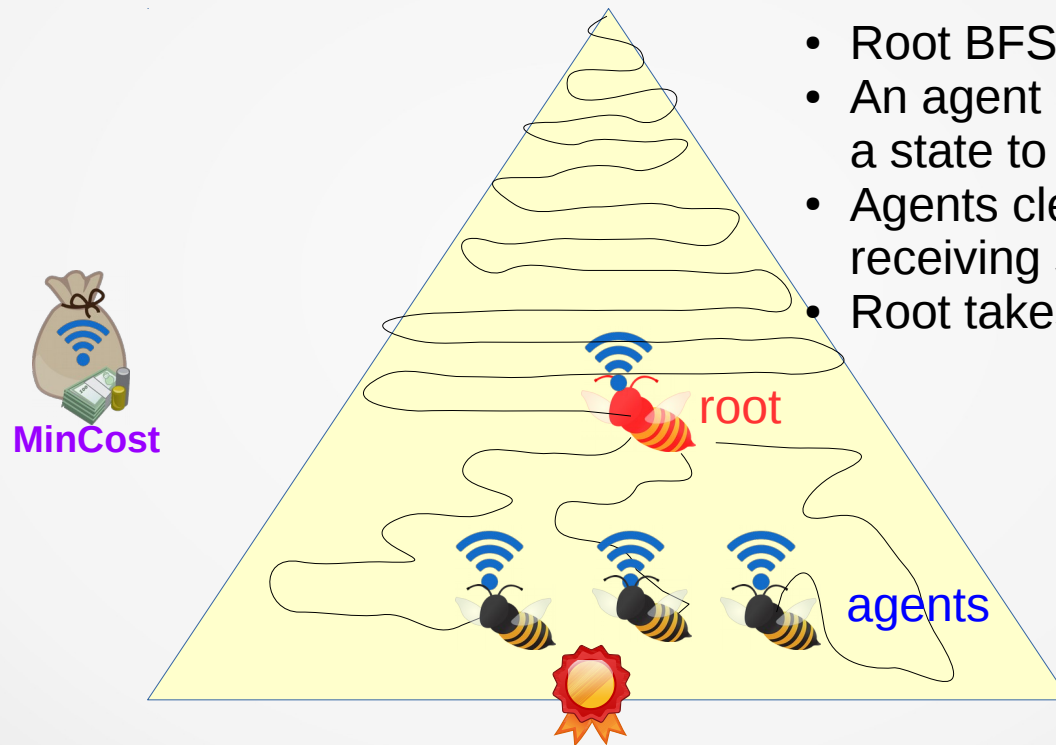


- One BFS, others RDFS.
- Share cost.

Algorithm in Appendix C

Swarm Time Optimal Reachability




- S-Agent: Cooperative Swarm Agent



- Root BFS, agents RDFS;
- An agent periodically asks root for a state to search from.
- Agents clear passed/waiting on receiving states. **Light weight.**
- Root takes charge of termination.

Algorithm in Appendix D

Swarm Time Optimal Reachability



Contents

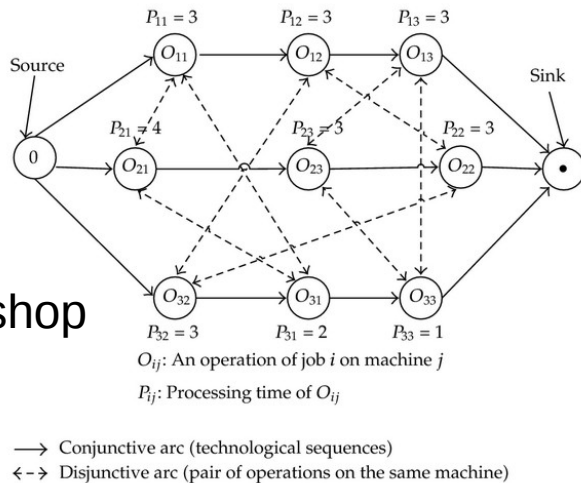
- Motivation
- Sequential Time Optimal Reachability
- Swarm Time Optimal Reachability
- **Experiments**
- Conclusion
- Future Work

Experiments

- Four timed scheduling models:
 - Job-shop-6 (jb-6): 6 people share a 4-section newspaper. Each person has own reading speed and favorable reading order on sections.
 - Aircraft-landing-15 (alp-15): 15 planes with different target landing time, need to land on 2 runways. Planes also have different landing intervals in between to avoid turbulence
 - Viking-bridge-15 (vik-15): 15 vikings with different walking speeds and sharing a single torch, cross a damaged bridge (Max 2 per) in darkness.
 - Task-graph-88 (task-88): non-preemptive schedule for 88 tasks with different precedence constraints on 4 CPUs with different frequencies (speeds).

Experiments

Jobshop




Airplane Landing



Viking Bridge



Task Graph

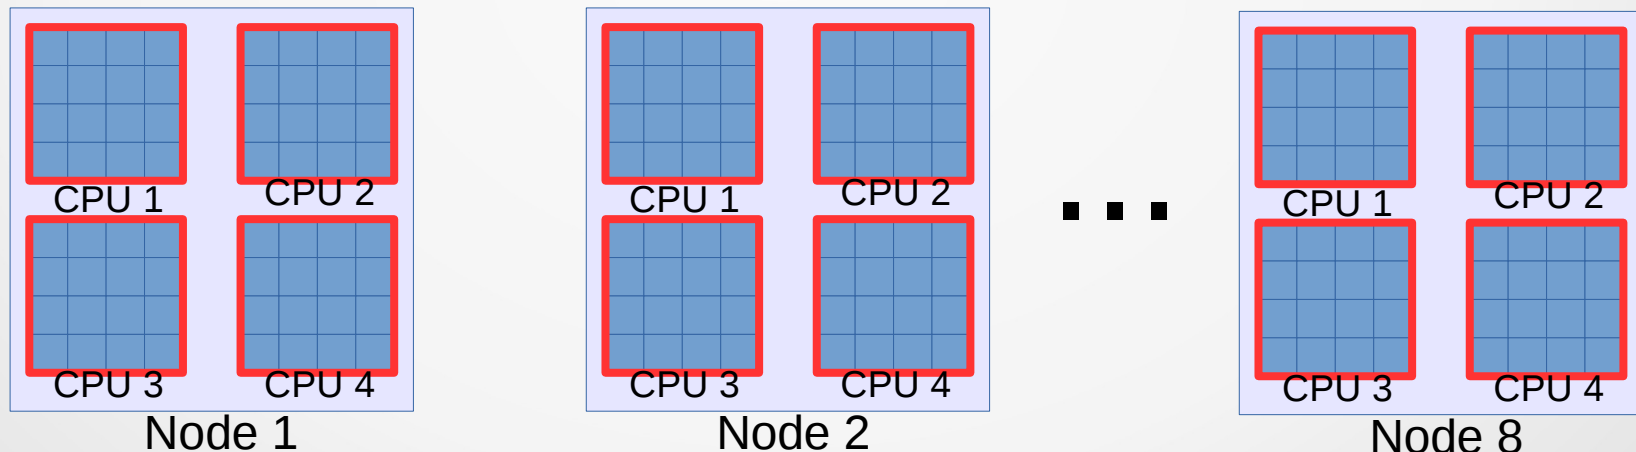


Experiments - Metric

- Metric of performance:
 - Metric 1: Shortest time to find optimal result t_{opt}
 - Metric 2: Time to terminate thus prove optimal result t_{prov}
 - Metric 3: Progressively improved near optimal results as a function of time, converge speed.
 - Metric 4: Memory consumption.

Experiments - Settings

- Experiment settings:
 - Cluster: NUMA architecture, 9 computing nodes, 1 TB per node, 4 CPU @ 2.3GHz per node, 16 cores per CPU, 574 cores total.
 - Algorithms: sequential-BFS/DFS/RDFS, P-RDFS, S-RDFS, S-Mix, S-Agent, 4-hour limit, 15 runs per core setting.



Jobshop-6 t_{opt} and t_{prov} (Metric 1&2)

| | Job-Shop-6 | | | | | | |
|----|------------|------------|-----------|------------|-----------|------------|--|
| #C | BFS | | DFS | | RDFS | | |
| | t_{opt} | t_{prov} | t_{opt} | t_{prov} | t_{opt} | t_{prov} | |
| 1 | 108 | 108 | 7877 | 9899 | 650 | 2283 | |

- BFS runs much faster than DFS/RDFS

Jobshop-6 t_{opt} and t_{prov} (Metric 1&2)

| Job-Shop-6 | | | | | | | |
|------------|-----------|------------|-----------|------------|-----------|------------|--|
| #C | BFS | | DFS | | RDFS | | |
| | t_{opt} | t_{prov} | t_{opt} | t_{prov} | t_{opt} | t_{prov} | |
| 1 | 108 | 108 | 7877 | 9899 | 650 | 2283 | |

| #C | P-RDFS | |
|-----|-----------|------------|
| | t_{opt} | t_{prov} |
| 2 | 535 | 1996 |
| 4 | 310 | 1730 |
| 8 | 176 | 1720 |
| 16 | 100 | 1591 |
| 32 | 65 | 1488 |
| 64 | 32 | 1452 |
| 128 | 4 | 1412 |
| 256 | 2 | 1365 |
| 512 | <1 | 1337 |

P-RDFS:

- Good speed up for t_{opt} , outperforms BFS above 16 cores.
- Limited speed up for t_{prov} .

Table 1: Runtime (s)

Jobshop-6 t_{opt} and t_{prov} (Metric 1&2)

| Job-Shop-6 | | | | | | |
|------------|-----------|------------|-----------|------------|-----------|------------|
| #C | BFS | | DFS | | RDFS | |
| | t_{opt} | t_{prov} | t_{opt} | t_{prov} | t_{opt} | t_{prov} |
| 1 | 108 | 108 | 7877 | 9899 | 650 | 2283 |

| #C | P-RDFS | | S-RDFS | |
|-----|-----------|------------|-----------|------------|
| | t_{opt} | t_{prov} | t_{opt} | t_{prov} |
| 2 | 535 | 1996 | 459 | 1821 |
| 4 | 310 | 1730 | 280 | 1683 |
| 8 | 176 | 1720 | 116 | 1653 |
| 16 | 100 | 1591 | 87 | 1506 |
| 32 | 65 | 1488 | 52 | 1494 |
| 64 | 32 | 1452 | 24 | 1385 |
| 128 | 4 | 1412 | 4 | 1367 |
| 256 | 2 | 1365 | <1 | 1341 |
| 512 | <1 | 1337 | <1 | 1263 |

S-RDFS vs. P-RDFS

- Sharing cost improves t_{opt} 15% and t_{prov} by 5%.

Table 1: Runtime (sec) of Job-Shop

Jobshop-6 t_{opt} and t_{prov} (Metric 1&2)

| Job-Shop-6 | | | | | | | |
|------------|-----------|------------|-----------|------------|-----------|------------|--|
| #C | BFS | | DFS | | RDFS | | |
| | t_{opt} | t_{prov} | t_{opt} | t_{prov} | t_{opt} | t_{prov} | |
| 1 | 108 | 108 | 7877 | 9899 | 650 | 2283 | |
| #C | P-RDFS | | S-RDFS | | S-Mix | | |
| | t_{opt} | t_{prov} | t_{opt} | t_{prov} | t_{opt} | t_{prov} | |
| 2 | 535 | 1996 | 459 | 1821 | 102 | 102 | |
| 4 | 310 | 1730 | 280 | 1683 | 102 | 102 | |
| 8 | 176 | 1720 | 116 | 1653 | 101 | 102 | |
| 16 | 100 | 1591 | 87 | 1506 | 99 | 101 | |
| 32 | 65 | 1488 | 52 | 1494 | 64 | 99 | |
| 64 | 32 | 1452 | 24 | 1385 | 24 | 98 | |
| 128 | 4 | 1412 | 4 | 1367 | 4 | 98 | |
| 256 | 2 | 1365 | <1 | 1341 | <1 | 98 | |
| 512 | <1 | 1337 | <1 | 1263 | <1 | 100 | |

S-Mix

- BFS instance allows best t_{prov} . and best t_{opt} between core 2 ~ 8.

Table 1: Runtime (sec) of Job-Shop-6 and Aircraft

Jobshop-6 t_{opt} and t_{prov} (Metric 1&2)

| | Job-Shop-6 | | | | | | | |
|-----|------------|------------|-----------|------------|-----------|------------|-----------|------------|
| #C | BFS | | DFS | | RDFS | | | |
| | t_{opt} | t_{prov} | t_{opt} | t_{prov} | t_{opt} | t_{prov} | | |
| 1 | 108 | 108 | 7877 | 9899 | 650 | 2283 | | |
| #C | P-RDFS | | S-RDFS | | S-Mix | | S-Agent | |
| | t_{opt} | t_{prov} | t_{opt} | t_{prov} | t_{opt} | t_{prov} | t_{opt} | t_{prov} |
| 2 | 535 | 1996 | 459 | 1821 | 102 | 102 | 39 | 99 |
| 4 | 310 | 1730 | 280 | 1683 | 102 | 102 | 25 | 98 |
| 8 | 176 | 1720 | 116 | 1653 | 101 | 102 | 17 | 97 |
| 16 | 100 | 1591 | 87 | 1506 | 99 | 101 | 14 | 97 |
| 32 | 65 | 1488 | 52 | 1494 | 64 | 99 | 13 | 96 |
| 64 | 32 | 1452 | 24 | 1385 | 24 | 98 | 16 | 95 |
| 128 | 4 | 1412 | 4 | 1367 | 4 | 98 | 29 | 95 |
| 256 | 2 | 1365 | <1 | 1341 | <1 | 98 | 28 | 95 |
| 512 | <1 | 1337 | <1 | 1263 | <1 | 100 | 31 | 96 |

- S-Agent vs. S-Mix
- Even better t_{opt} from core 2 ~ 64
 - Equally good t_{prov} .

Table 1: Runtime (sec) of Job-Shop-6 and Aircraft-Landing-15

Aircraft-15 t_{opt} and t_{prov} (Metric 1&2)

| Aircraft-Landing-15 | | | | | | | | |
|---------------------|-----------|------------|-----------|------------|-----------|------------|-----------|------------|
| #C | BFS | | DFS | | RDFS | | | |
| | t_{opt} | t_{prov} | t_{opt} | t_{prov} | t_{opt} | t_{prov} | | |
| 1 | 157 | 159 | 73 | 428 | 191 | 964 | | |
| #C | P-RDFS | | S-RDFS | | S-Mix | | S-Agent | |
| | t_{opt} | t_{prov} | t_{opt} | t_{prov} | t_{opt} | t_{prov} | t_{opt} | t_{prov} |
| 2 | 87 | 845 | 93 | 873 | 91 | 91 | 66 | 90 |
| 4 | 42 | 737 | 35 | 729 | 46 | 90 | 21 | 91 |
| 8 | 22 | 712 | 16 | 727 | 17 | 89 | 25 | 88 |
| 16 | 17 | 676 | 11 | 699 | 11 | 89 | 6 | 86 |
| 32 | 2 | 638 | 1 | 633 | 2 | 89 | 12 | 86 |
| 64 | <1 | 604 | <1 | 582 | <1 | 88 | 5 | 85 |
| 128 | <1 | 597 | <1 | 620 | <1 | 88 | 2 | 84 |
| 256 | <1 | 581 | <0.1 | 565 | <0.1 | 88 | 7 | 84 |
| 512 | <1 | 585 | <0.1 | 573 | <0.1 | 117 | 5 | 86 |

- Similar pattern as Jobshop
- Sharing costs allow S-Mix and S-Agent 45% faster in t_{prov} than BFS.

Table 1: Runtime (sec) of Job-Shop-6 and Aircraft-Landing-15

Viking-15 t_{opt} and t_{prov} (Metric 1&2)


mincost=669930, timestamp=9.5024
 mincost=669925, timestamp=10.455
 mincost=669910, timestamp=10.467
 mincost=669905, timestamp=11.420
 mincost=669895, timestamp=12.481
 mincost=669890, timestamp=13.408
 mincost=669885, timestamp=13.418
 mincost=669880, timestamp=14.344
 mincost=669870, timestamp=14.350
 mincost=669865, timestamp=14.361
 mincost=669815, timestamp=15.330
 mincost=669795, timestamp=16.324
 mincost=669790, timestamp=17.245
 mincost=669785, timestamp=18.332
 mincost=669780, timestamp=19.289
 mincost=669775, timestamp=20.205
 mincost=669770, timestamp=21.323
 mincost=669765, timestamp=21.336
 mincost=669760, timestamp=22.363
 mincost=669755, timestamp=23.332
 mincost=669750, timestamp=24.316
 mincost=669745, timestamp=25.404

... Fine grain cost
 mincost=220

| #C | BFS | | DFS | | RDFS | | | |
|-----------------------|-----------|------------|-----------|------------|-----------|------------|-----------|------------|
| | t_{opt} | t_{prov} | t_{opt} | t_{prov} | t_{opt} | t_{prov} | | |
| 1 | 70 | 70 | - | - | - | - | | |
| #C | P-RDFS | | S-RDFS | | S-Mix | | S-Agent | |
| | t_{opt} | t_{prov} | t_{opt} | t_{prov} | t_{opt} | t_{prov} | t_{opt} | t_{prov} |
| 2 | - | - | - | - | 96 | 96 | 75 | 75 |
| 4 | - | - | - | - | 87 | 87 | 72 | 72 |
| 8 | - | - | - | - | 96 | 96 | 72 | 72 |
| 16 | - | - | - | - | 92 | 92 | 83 | 83 |
| 32 | - | - | - | - | 100 | 100 | 94 | 94 |
| 64 | - | - | - | - | 107 | 107 | 77 | 77 |
| 128 | - | - | - | - | 112 | 112 | 74 | 74 |
| 256 | - | - | - | - | 186 | 186 | 71 | 71 |
| 512 | - | - | - | - | 215 | 215 | 73 | 73 |
| “-”: denotes timeout. | | | | | | | | |

Table 2: Runtime (sec) of Viking-Bridge-15

- Sharing costs decelerate S-Mix and S-Agent.



Experiments – t_{opt} and t_{prov} (Metric 1&2)

- Summary for (Metric 1&2)
 - Exchanging cost can generally improve t_{opt} , at low core settings (2 ~ 64);
 - S-Mix and S-Agent combines benefits of BFS and RDFS. They can report results and terminate faster;
 - Costs reported by RDFS instances can help BFS root in pruning, but fine-grained improved costs may backfire.

Jobshop-6 Results versus time (Metric 3)




Figure 1: Cost vs. Runtime for Job-Shop-6

Task-88 Results versus time (Metric 3)

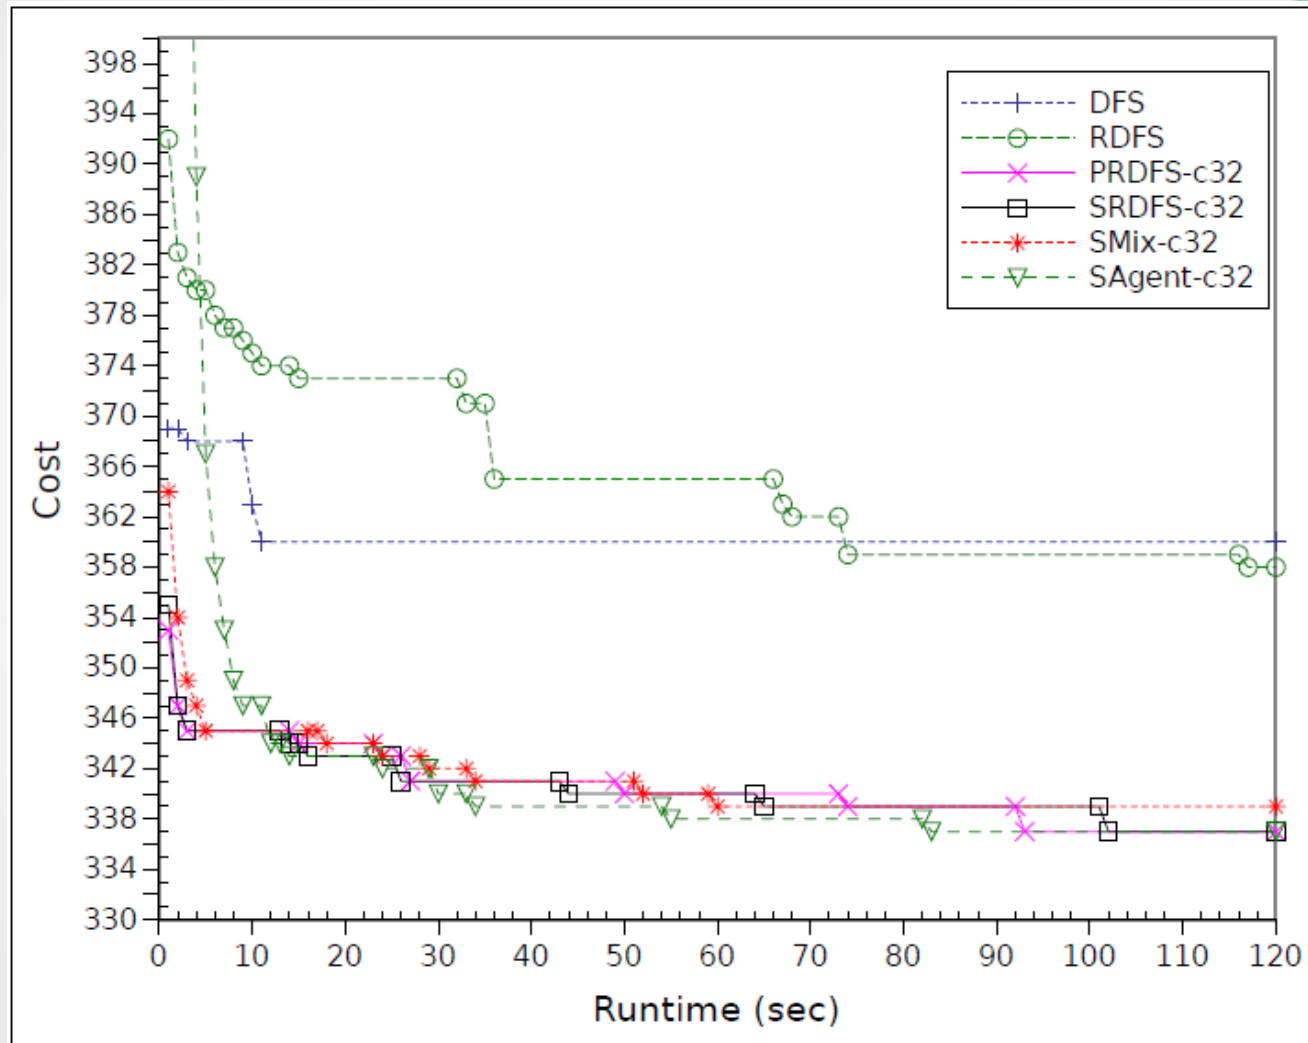


Figure 2: Cost vs. Runtime for Task-Graph-88

Aircraft-15 Results versus time (Metric 3)




Figure 3: Cost vs. Runtime for Aircraft-Landing-15

Viking-15 Results versus time (Metric 3)

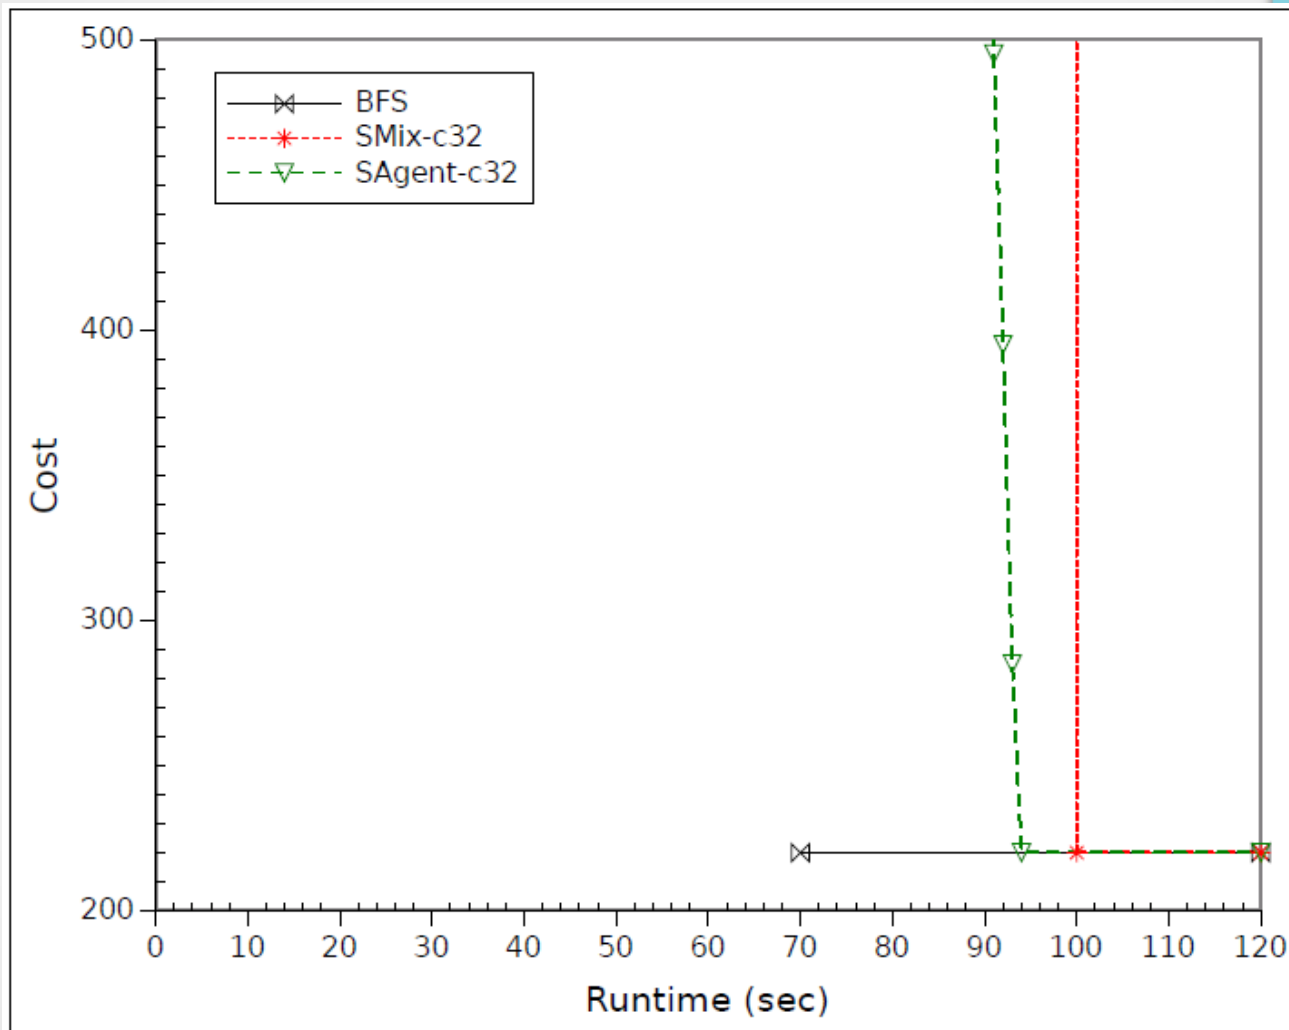


Figure 4: Cost vs. Runtime for Viking-Bridge-15

Experiments – Results versus time (Metric 3)

- Summary (Metric 3)
 - Swarm algorithms are generally faster than sequential algorithms at finding near optimal results, and they work equally well.

Experiments – Memory consumption (Metric 4)

| Models | jb-6 | alp-15 | vik-15 | task-88 |
|--|------|--------|--------------|-------------|
| BFS | 152 | 124 | 408 | <i>5068</i> |
| DFS | 1754 | 127 | <i>13297</i> | <i>3060</i> |
| RDFS | 606 | 191 | <i>12790</i> | <i>3668</i> |
| P-RDFS-c32 | 593 | 187 | <i>13555</i> | <i>3849</i> |
| S-RDFS-c32 | 591 | 190 | <i>15303</i> | <i>3873</i> |
| S-Mix-c32 | 63 | 51 | 230 | <i>3914</i> |
| S-Agent-c32 | 21 | 13 | 100 | <i>1277</i> |
| <i>Italic font denotes termination due to timeout.</i> | | | | |

RDFS instances/agents
amortize average

Table 3: Resident Memory (MB) per UPPAAL Instance

- Summary (Metric 4)
 - S-Agent has the best average memory footprint.

Contents

- Motivation
- Sequential Time Optimal Reachability
- Swarm Time Optimal Reachability
- Experiments
- Conclusion
- Future Work

Conclusion

- Swarm algorithms generally find optimal (or near optimal) results faster than sequential algorithms.
- S-Mix and S-Agent combine the benefits of BFS and RDFS. They find results and terminate fast.
- S-Agent has smaller memory footprint because agents do not keep state space.
- Sharing cost is useful for speed up at lower core settings.

Future Work

- Swarm intelligence algorithms from AI:
 - Ant colony optimization algorithms (ACO)
 - Particle swarm optimization (PSO)
 - Rapidly-exploring Random Tree (RRT)
- Parallel- and distributed algorithms
- Extend to more general priced time automata (PTA).

Thank You for Your Attention



Appendix A: P-RDFS Algorithm

- P-RDFS

Algorithm 2: Basic Swarm RDFS Time Optimal Reachability (P-RDFS)

PASSED $\leftarrow \emptyset$, WAITING $\leftarrow \{(\ell_0, Z_0)\}$, COST $\leftarrow \infty$

ORDER \leftarrow RDFS, TERMINATE \leftarrow FALSE

Procedure Main()

```
1  while WAITING  $\neq \emptyset$  and  $\neg$ TERMINATE do
2    Search()
3    stop other instances
4  return COST
```

All do RDFS

First-complete-terminate

Appendix B: S-RDFS Algorithm

- S-RDFS:

Algorithm 3: Cooperative Swarm RDFS Time Optimal Reachability (S-RDFS)

(Local Variables)

PASSED $\leftarrow \emptyset$, WAITING $\leftarrow \{(\ell_0, Z_0)\}$

COST $\leftarrow \infty$, ECOST $\leftarrow \infty$

Cost received externally

ORDER \leftarrow RDFS, TERMINATE \leftarrow FALSE

(Message Types)

Notify better cost

UPDATE

Procedure Main()

1 while WAITING $\neq \emptyset$ and \neg TERMINATE do

// lines [2,3] are called atomically

2 if ECOST < COST then COST \leftarrow ECOST

3 if COST < ECOST then

| Update(COST), Broadcast(UPDATE, COST)

4 Search()

5 stop other instances

6 return COST

Procedure Update(NEWCost) // called atomically

7 if NEWCost < ECOST then ECOST \leftarrow NEWCost

(Message Processing Rules)

8 When a node receives UPDATE(Ncost), Update(Ncost).

Before local search
Compare costs

Message handler

Appendix C: S-Mix Algorithm

- S-Mix

Algorithm 4: Cooperative Swarm Mix Time Optimal Reachability (S-Mix)

(Local Variables)

$\text{ORDER} \leftarrow \begin{cases} \text{BFS} & \text{if } p = 0, \\ \text{RDFS} & \text{otherwise.} \end{cases}$ One (BFS), others (RDFS)

The rest is the same as Algorithm 3

Appendix D: S-Agent Algorithm

- S-Agent

Algorithm 5: Cooperative Swarm Agent Time Optimal Reachability (S-Agent)

(Local Variables)

$\text{PASSED} \leftarrow \emptyset, \text{COST} \leftarrow \infty, \text{ECOST} \leftarrow \infty$

$\text{WAITING} \leftarrow \begin{cases} \{(\ell_0, Z_0)\} & \text{if } p = \text{root}, \\ \emptyset & \text{otherwise.} \end{cases}$

$\text{ORDER} \leftarrow \begin{cases} \text{BFS} & \text{if } p = \text{root}, \\ \text{RDFS} & \text{otherwise.} \end{cases}$

Root (BFS), agents (RDFS)
Root searches from initial state
Agents have empty waiting list

(Message Types)

UPDATE, REQUEST, ISSUE

Two more messages

Procedure MainRoot()

Root

```
1  while WAITING  $\neq \emptyset$  do
   |  same as lines [2-4] in Algorithm 3
2  stop all agents
3  return COST
```

Root stops agents

Appendix D: S-Agent Algorithm

- S-Agent Cont.

Agent

```
Procedure MainAgent()  
4  ITERATION  $\leftarrow$  0, LIMIT  $\in \mathbb{N}$ , TERMINATE  $\leftarrow$  FALSE  
5  while  $\neg$ TERMINATE do  
6    if WAITING =  $\emptyset$  then  
      | Send(REQUEST, Root), Recv(ISSUE)  
      // lines [7,8] are called atomically  
7    if ECOST < Cost then Cost  $\leftarrow$  ECOST  
8    if Cost < ECOST then  
      | Update(Cost), Broadcast(UPDATE, Cost)  
9    Search(), ITERATION  $\leftarrow$  ITERATION + 1  
10   if ITERATION  $\geq$  LIMIT then  
11     | Send(REQUEST, Root), Recv(ISSUE),  
      | ITERATION  $\leftarrow$  0  
  
(Message Processing Rules)  
12 When a node receives UPDATE(NCOST): Update(NCOST).  
13 When root receives REQUEST from agent  $i$ : select  $(\ell', Z')$   
   from WAITING, Send(ISSUE,  $(\ell', Z'), i$ ).  
14 When an agent receives ISSUE $\langle(\ell^*, Z^*)\rangle$  from root:  
   WAITING  $\leftarrow \emptyset$ , PASSED  $\leftarrow \emptyset$ , WAITING  $\leftarrow \{(\ell^*, Z^*)\}$ .
```

Send request to root,
then blocking receive a
state from root

Send request in a period


Root hands Request
Agents handle Issue

Agent is light weight

Agents clear old state
space when received
a new state

Appendix E: Scalability problem

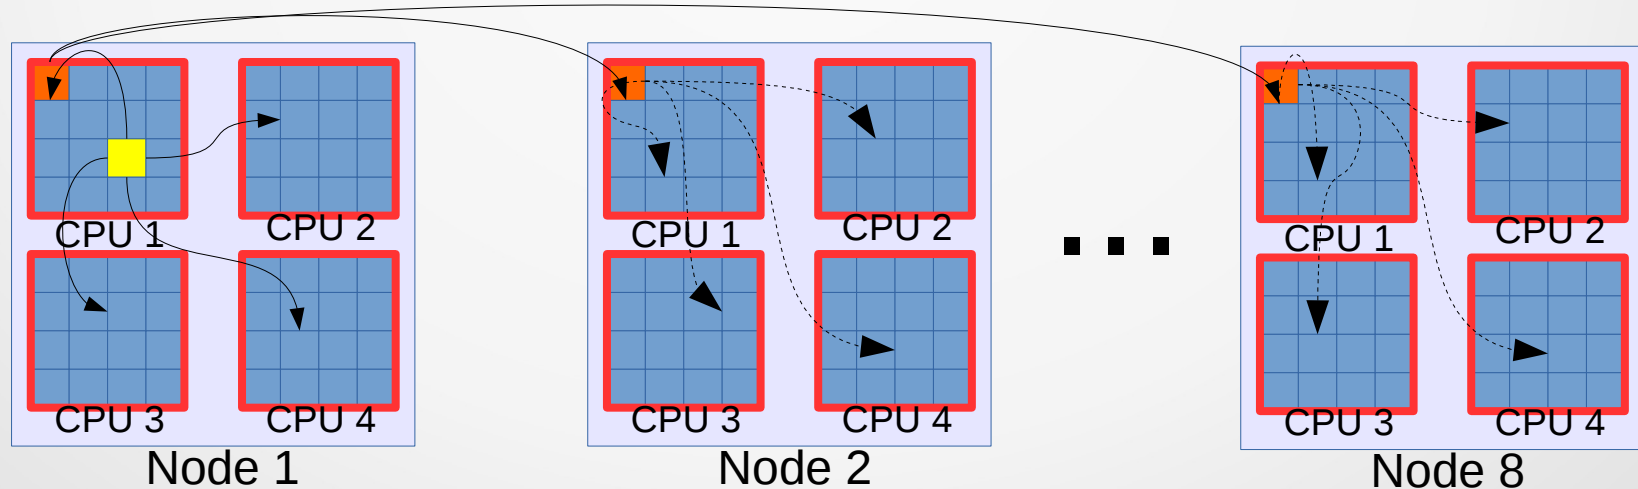
- Scalability problem:
 - Message **flood**: 512 instances broadcast costs at the same time will overload the cluster severely.
 - Cluster: NUMA architecture, 9 computing nodes, 1 TB per node, 4 CPU @ 2.3GHz per node, 16 cores per CPU, 574 cores total.



7×64 (=448) inter-node messages overload cluster heavily!!!

Appendix E: Scalability problem

- Optimization for scalability
 - Overcome by using cluster architecture, virtual topology;
 - Virtual node: one master core, other normal core;
 - Normal core broadcasts within same virtual node;
 - Only master core send/relay inter-node message.



Now only 7 inter-node messages. Scalability is retained!